



mp

max planck institut
informatik

SIC Saarland Informatics
Campus

High Level Computer Vision - Introduction

@ April 9, 2025

Bernt Schiele

<https://cms.sic.saarland/hlcvss25/>

**Max Planck Institute for Informatics & Saarland University,
Saarland Informatics Campus Saarbrücken**



mp

max planck institut
informatik

SIC Saarland Informatics
Campus

Computer Vision and Machine Learning Group @ Max-Planck-Institute for Informatics



Bernt Schiele
Computer Vision &
Machine Learning

Jonas Fischer
Explainable Machine
Learning



Jan Eric Lenssen
Geometric Representation
Learning

Margret Keuper
Robust Visual Learning
U Mannheim



High Level Computer Vision

- Lecturer:
 - ▶ Bernt Schiele (schiele@mpi-inf.mpg.de)
- Assistant:
 - ▶ Amin Parchami-Araghi (mparcham@mpi-inf.mpg.de)
 - ▶ Nhi Pham (nhipham@mpi-inf.mpg.de)
 - ▶ Sophia Wiedmann
 - ▶ Mostafa Abdelgawad
- Language:
 - ▶ English
- mailing list for announcements etc.
 - ▶ register online: <https://cms.sic.saarland/hcvss25/>



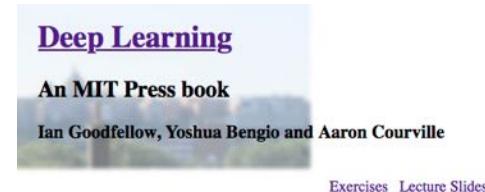
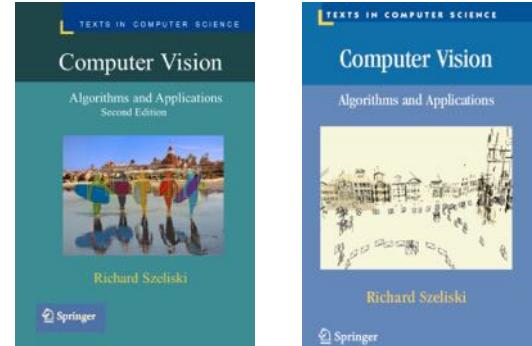
Lecture & Tutorial

- Officially: 2V (lecture) + 2Ü (exercise/tutorial) - SWS 002
 - ▶ Lecture: Wed: 10:15am - 12pm
 - ▶ Tutorial: Mon: 10:15am - 12pm
- typically 1 exercise sheet every 1-2 weeks
 - ▶ working in **groups of 2 students**
 - ▶ some pencil and paper, mostly practical including a larger project
 - ▶ larger project: we/you propose projects, mentoring, final presentation
 - ▶ Mon, April 14th: Python tutorial
- Exam
 - ▶ written exam
 - ▶ exercise & project participation **mandatory** to register for exam
 - ▶ up to 1 **grade improvement** (50% project, 50% exercises)



Material

- For "non-deep-learning" parts of the lecture:
 - ▶ available online
<http://szeliski.org/Book>
- Background on deep learning:
Deep Learning Book
 - ▶ available online
<http://deeplearningbook.org>



The Deep Learning textbook is a resource intended to help students and practitioners enter the field of machine learning in general and deep learning in particular. The online version of the book is now complete and will remain available online for free. The print version will be available for sale soon. For up to date announcements, join our [mailing list](#).

Citing the book

To cite this book, please use this bibtex entry:

```
@unpublished{Goodfellow-et-al-2016-Book,  
    title={Deep Learning},  
    author={Ian Goodfellow, Yoshua Bengio, and Aaron Courville},  
    note={Book in preparation for MIT Press},  
    url={(http://www.deeplearningbook.org)},  
    year={2016}  
}
```

Overview Lecture (preliminary)

- Convolutional Neural Networks (CNNs) for High-Level Computer Vision Tasks:
 - Image Classification, Object Detection (Recognition, 2D Localization), Semantic Segmentation
 - Data preprocessing & Batch normalization, CNN architectures
- Transformer Architectures for High-Level Vision Tasks
 - Vision Transformers for image classification,
 - DETR for object detection, Semantic Segmentation
- Generative Models
 - Generative Adversarial Networks (GANs), Variational Auto-Encoders (VAEs), (Stable) Diffusion, ...
- Self-Supervised Learning
 - How to learn with less or even no supervision
- Foundational Models
 - Vision-Language Models, CLIP, Segment Anything Model (SAM), ...
- Recent Trends and Works
 - B-Cos Networks (inherently interpretable), Flamingo, ConvNeXt, ...

Goals of today's lecture

- First intuitions about
 - ▶ What is computer vision?
 - ▶ What does it mean to see and how do we (as humans) do it?
 - ▶ How can we make this computational?
- Applications & Appetizers
- Case Study
 - ▶ Object Recognition — intuition from human vision...
- Image Classification
 - ▶ using data-driven approaches (machine learning)
 - ▶ K-nearest neighbor classifier



Why Study Computer Vision

- Science — How do Humans “See”
 - ▶ Foundations of perception. How do WE as humans see?
 - ▶ computer vision to explore “computational model of human vision”
- Engineering — How can We Make Machines “See”
 - ▶ How do we build systems that perceive the world
 - ▶ computer vision to solve real-world problems
(e.g. self-driving cars to detect pedestrians)
- Applications
 - ▶ medical imaging (computer vision to support medical diagnosis, visualization)
 - ▶ car-industry (lane-keeping, pre-crash intervention, ...)
 - ▶ security & surveillance (to follow/track people at the airport, train-station, ...)
 - ▶ entertainment (vision-based interfaces for games)
 - ▶ graphics (image-based rendering, vision to support realistic graphics)
 - ▶ ...



Some Applications

- License Plate Recognition
 - ▶ London Congestion Charge
 - ▶ http://en.wikipedia.org/wiki/London_congestion_charge
- Security & Surveillance
 - ▶ Face Recognition
 - ▶ Airport Security
(People Tracking)
- Medical Imaging
 - ▶ (Semi-)automatic segmentation and measurements
- Autonomous Driving & Robotics
 - ▶ Waymo, Chrysler, etc.

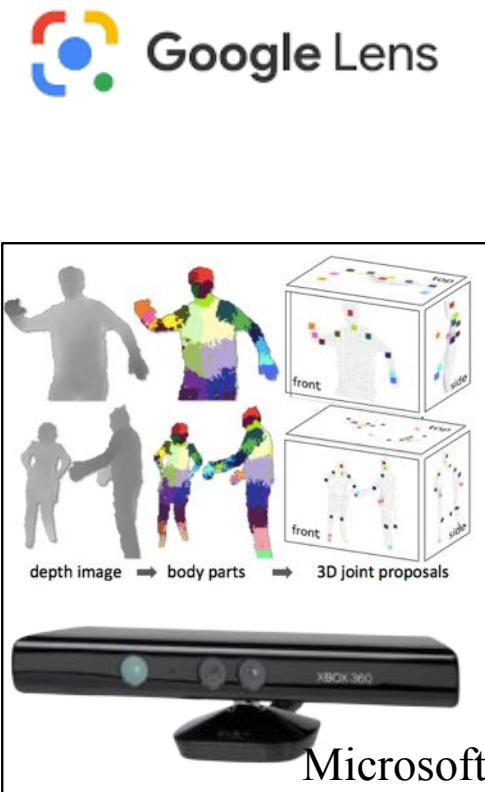


image: digital trends



More Applications

- Vision on Cellphones:
 - ▶ e.g. Google Lens, ...
- Vision for Interfaces:
 - ▶ e.g. Microsoft Kinect
- Reconstruction



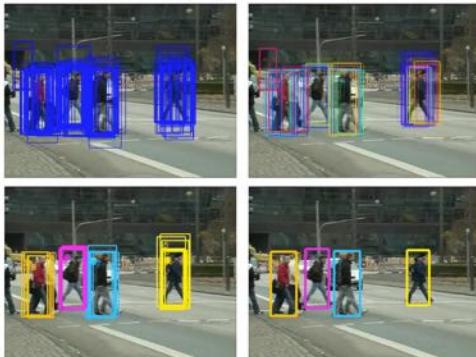


Applications & Appetizers

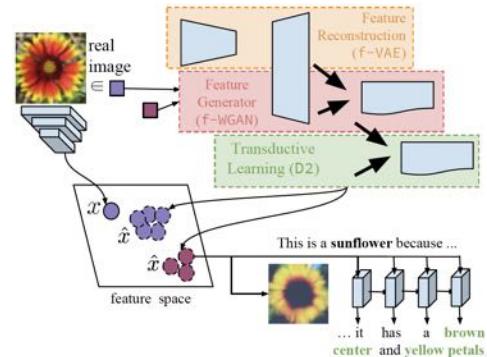
... work from our group

Some Research Highlights from the Past

Multi-Person Tracking



Zero- and Few-Shot Learning



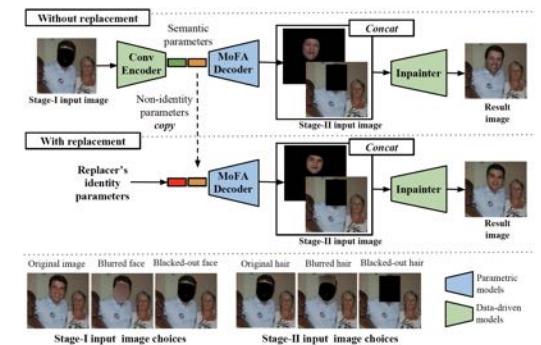
Multi-Person Human Pose Estimation & Tracking



Video Segmentation



Visual Privacy

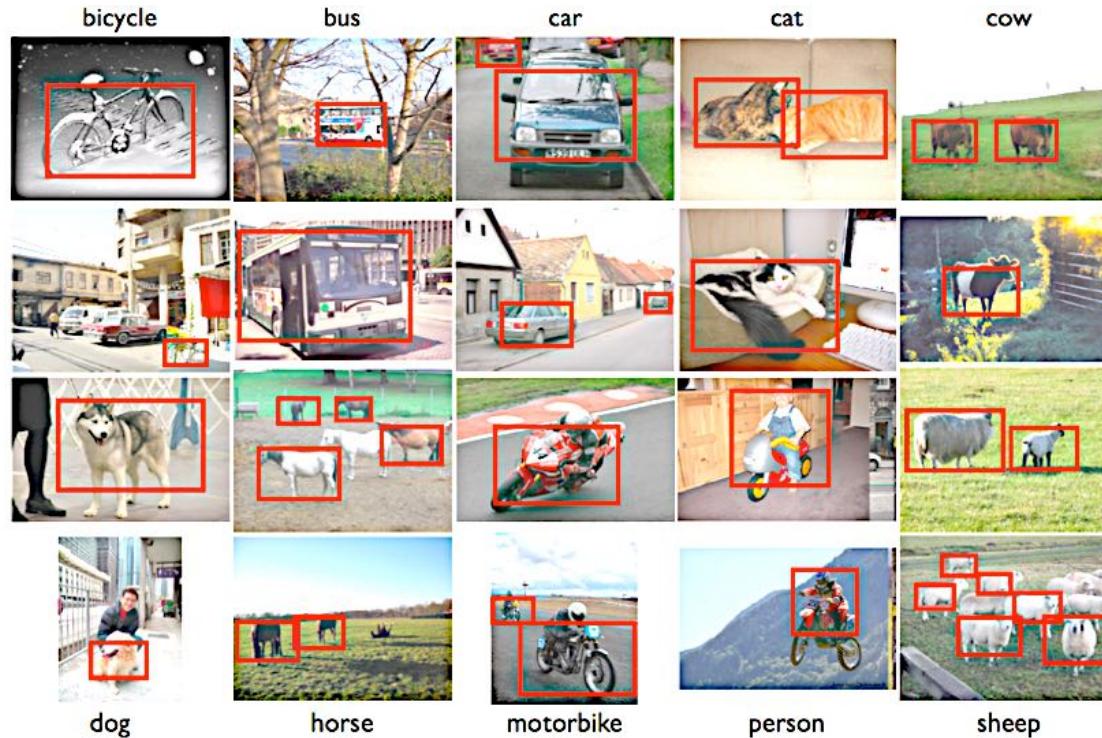


Datasets...

Caltech Pedestrians



Detection & Recognition of Visual Categories



- Challenges:
- multi-scale
 - multi-view
 - multi-class
 - varying illumination
 - occlusion
 - cluttered background
 - articulation
 - high intraclass variance
 - low interclass variance

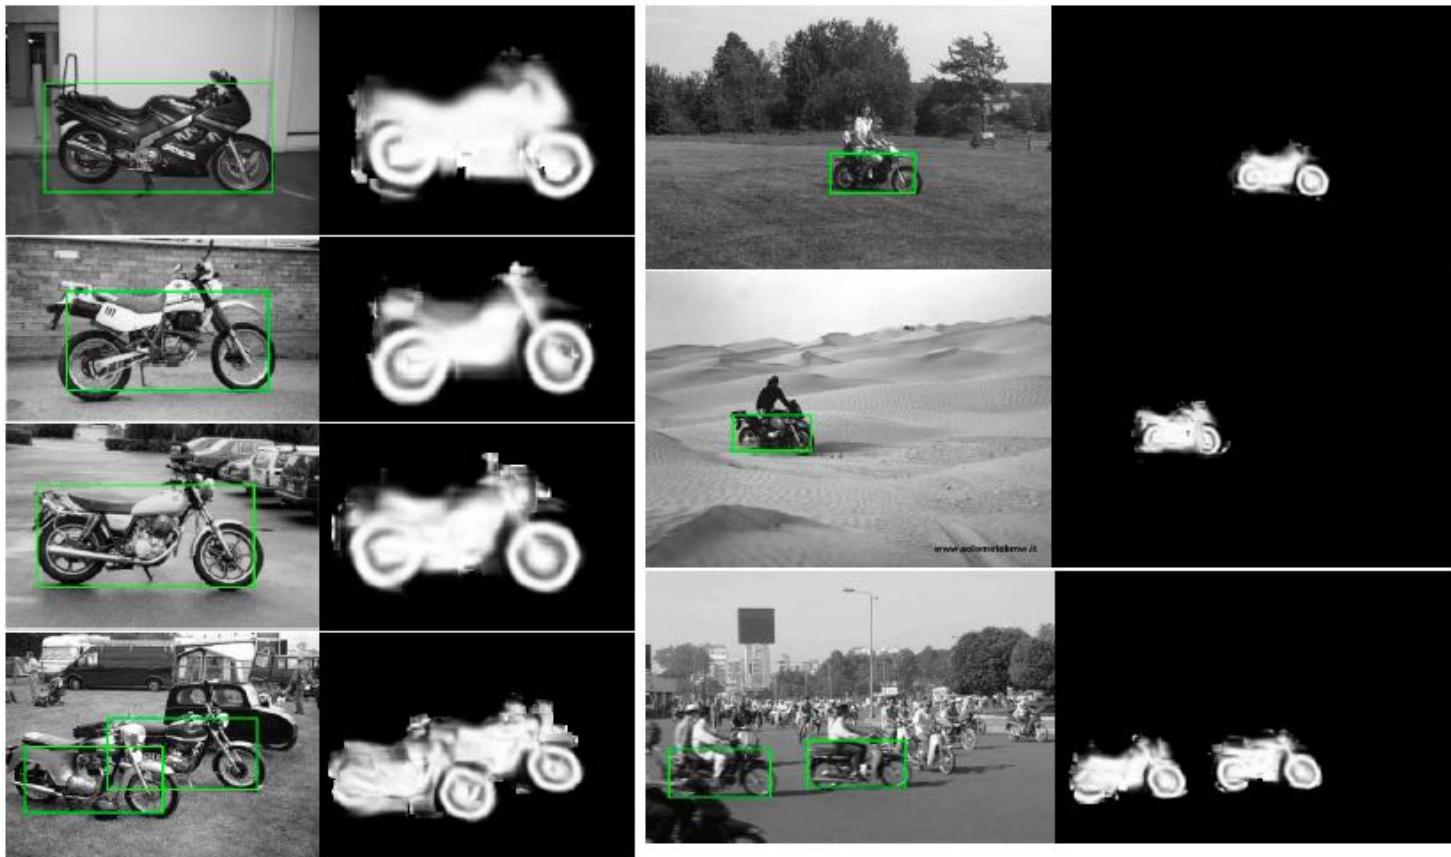
Challenges of Visual Categorization

- high intra-class variation
- low inter-class variation

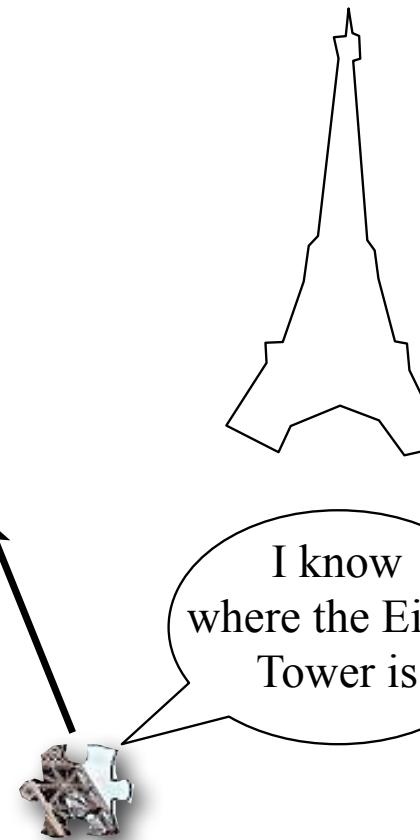
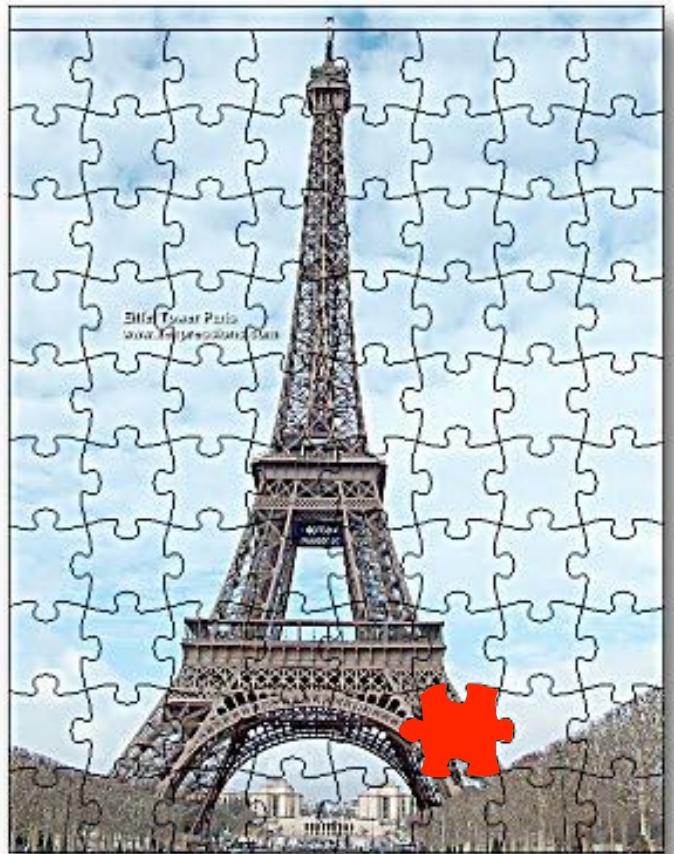


- high intra-class variation

Sample Category: Motorbikes



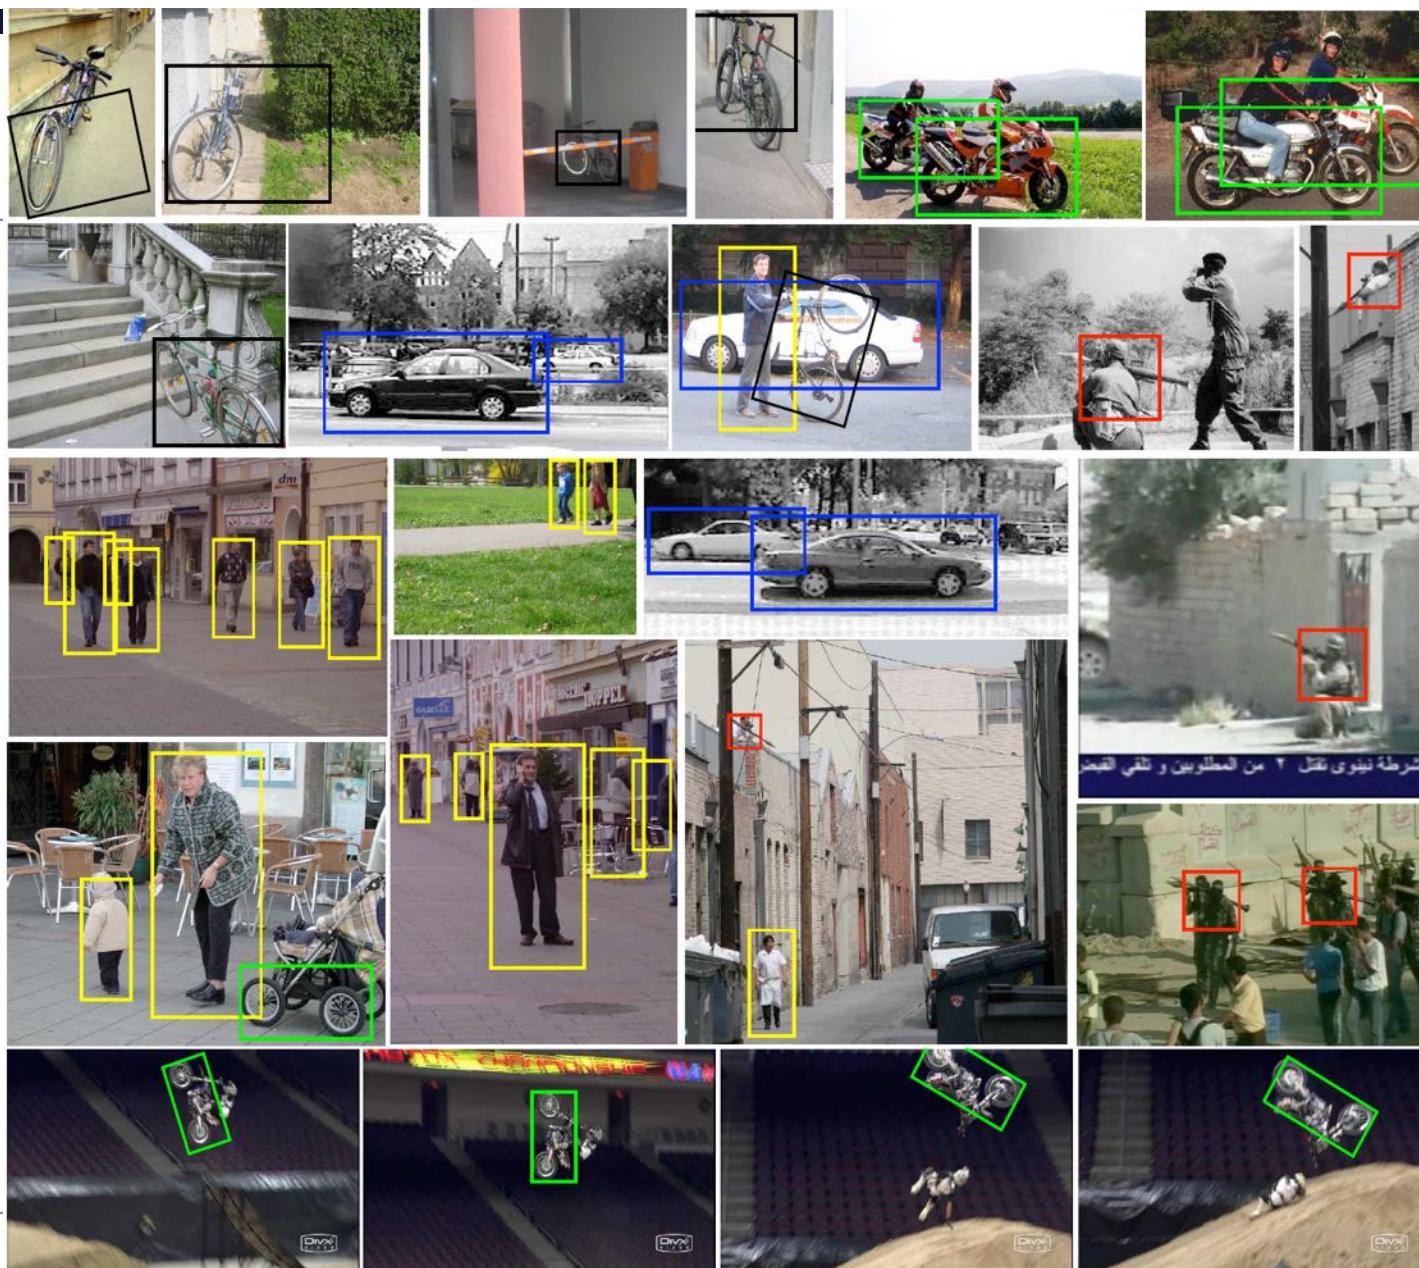
Basic Idea



global

I know
where the Eiffel
Tower is

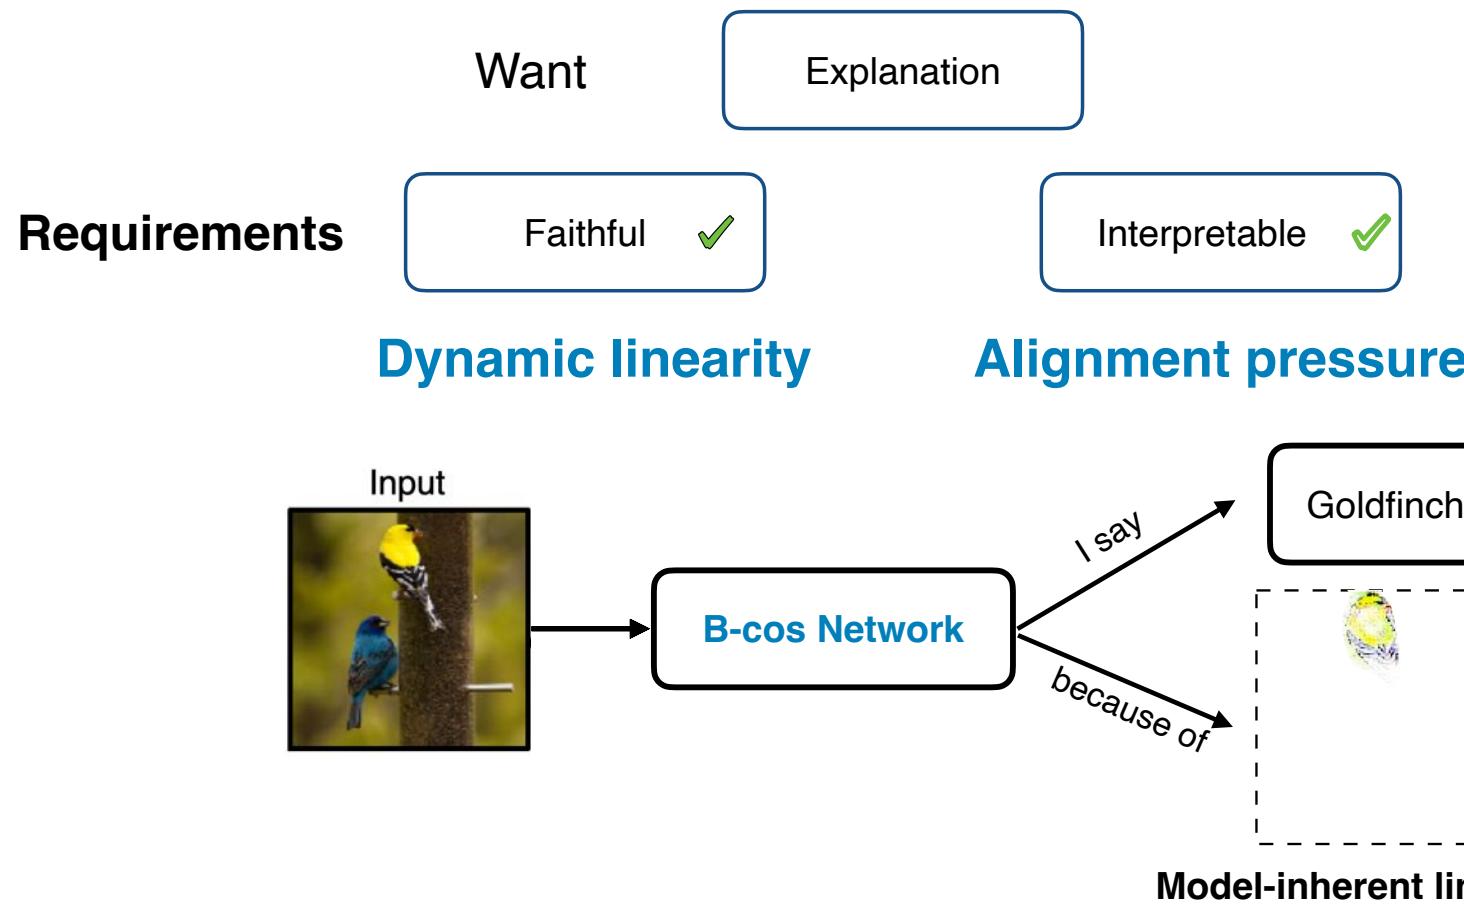
local



Computer Vision & Machine Learning: Some Research Themes

- **Interpretability, Robustness** and **Security** of Deep Learning in Computer Vision
 - ▶ **Inherently Interpretable** Deep Neural networks — CVPR'21, CVPR'22, PAMI'23, PAMI'24
 - ▶ **Robustness** of Deep Models:
Bright and Dark Side of Scene Context — NeurIPS'18, CVPR'19, ECCV'20
 - ▶ **Security** of Deep Models
Reverse Engineering and **Stealing** of Deep Models — ICLR'18, CVPR'19, ICLR'20

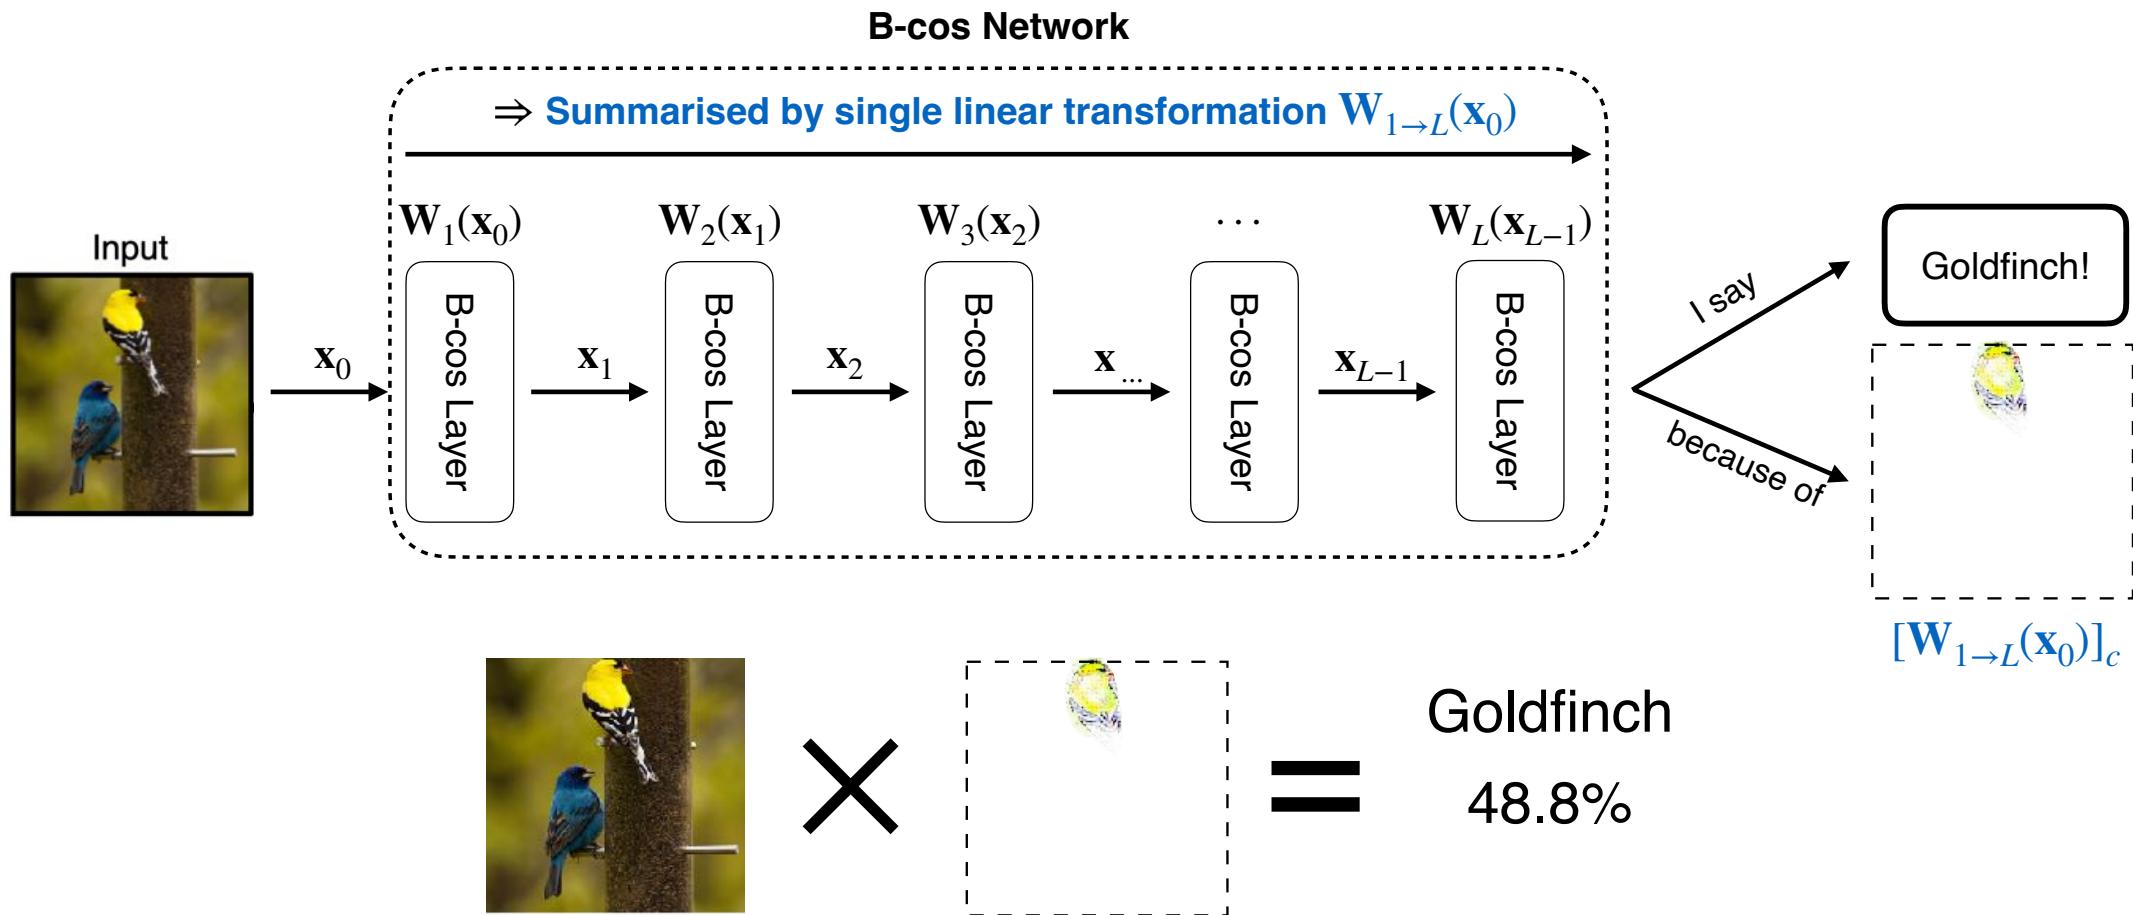
Motivation: we aim for Inherent Interpretability



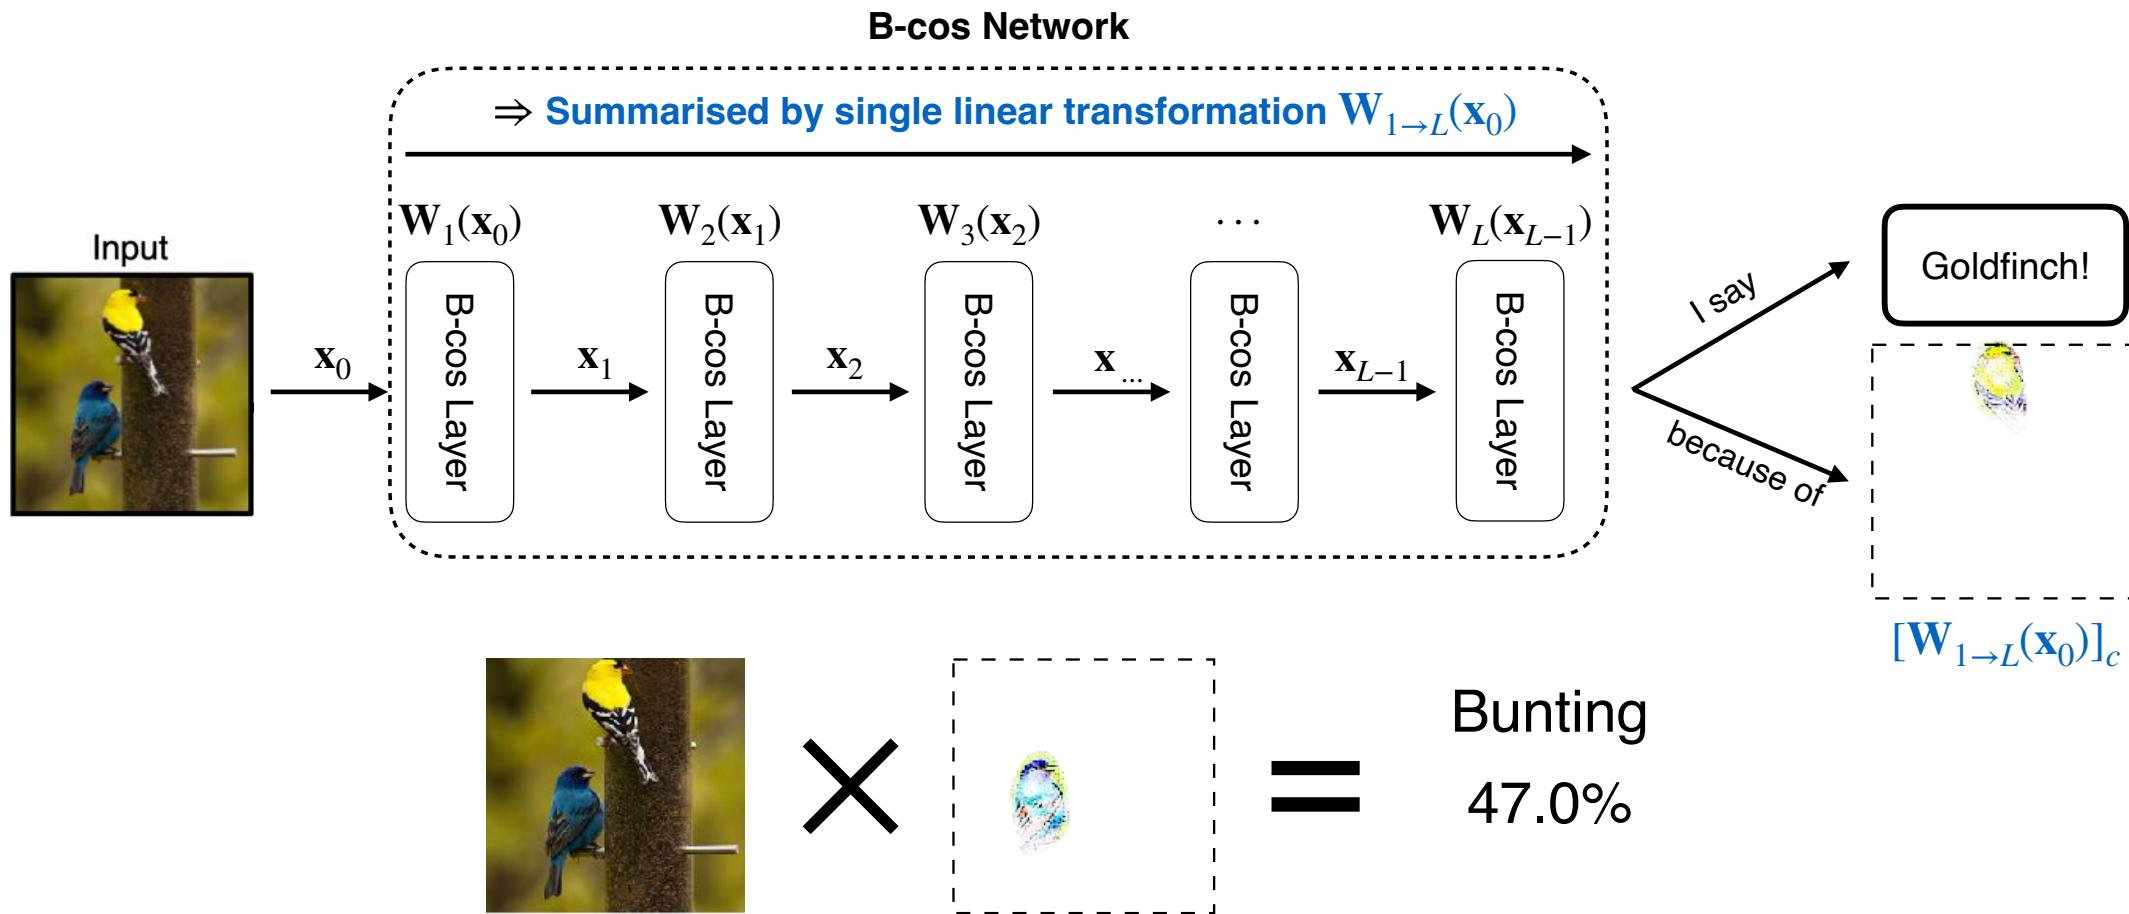
References: 'Requirements' (Gilpin et al., 2018), VGG-11 (Simonyan et al., 2014), Grad (Baehrens et al., 2010), Guided Backpropagation (Springenberg et al., 2014), Sanity check (Adebayo et al., 2018)



Dynamic linearity



Dynamic linearity



Visualisations of $W_{1 \rightarrow L}(x)$

Input image



Input image



Computer Vision & Machine Learning: Some Research Themes

- **Interpretability, Robustness** and **Security** of Deep Learning in Computer Vision
 - ▶ **Inherently Interpretable** Deep Neural networks — CVPR'21, CVPR'22
 - ▶ **Robustness** of Deep Models:
Bright and Dark Side of Scene Context — NeurIPS'18, CVPR'19, ECCV'20
 - ▶ **Security** of Deep Models
Reverse Engineering and **Stealing** of Deep Models — ICLR'18, CVPR'19, ICLR'20

Recognition: the Role of Context

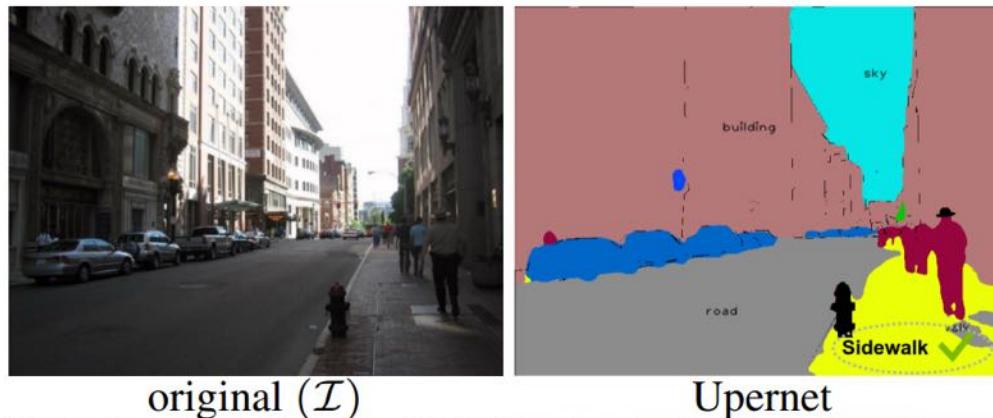
- Antonio Torralba



The Bright and the Dark Side of Scene Context

- Current models heavily rely on scene context:

- ▶ Original image with cars on the left side:



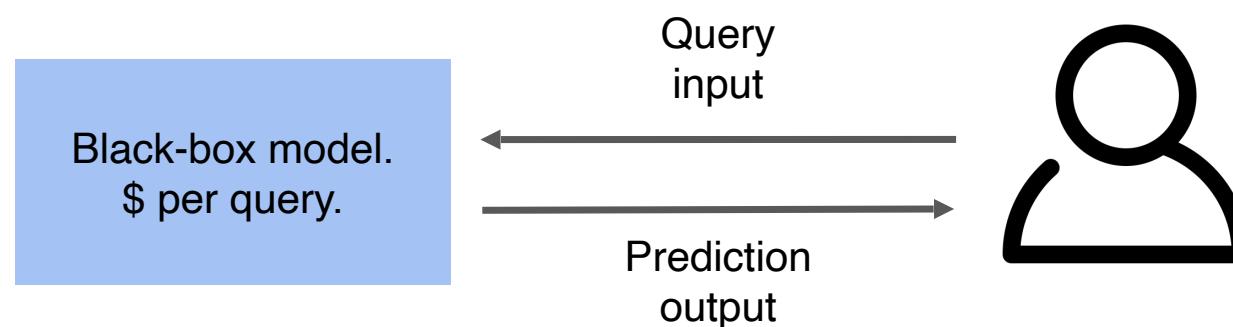
- ▶ Same image without those cars:

Computer Vision & Machine Learning: Some Research Themes

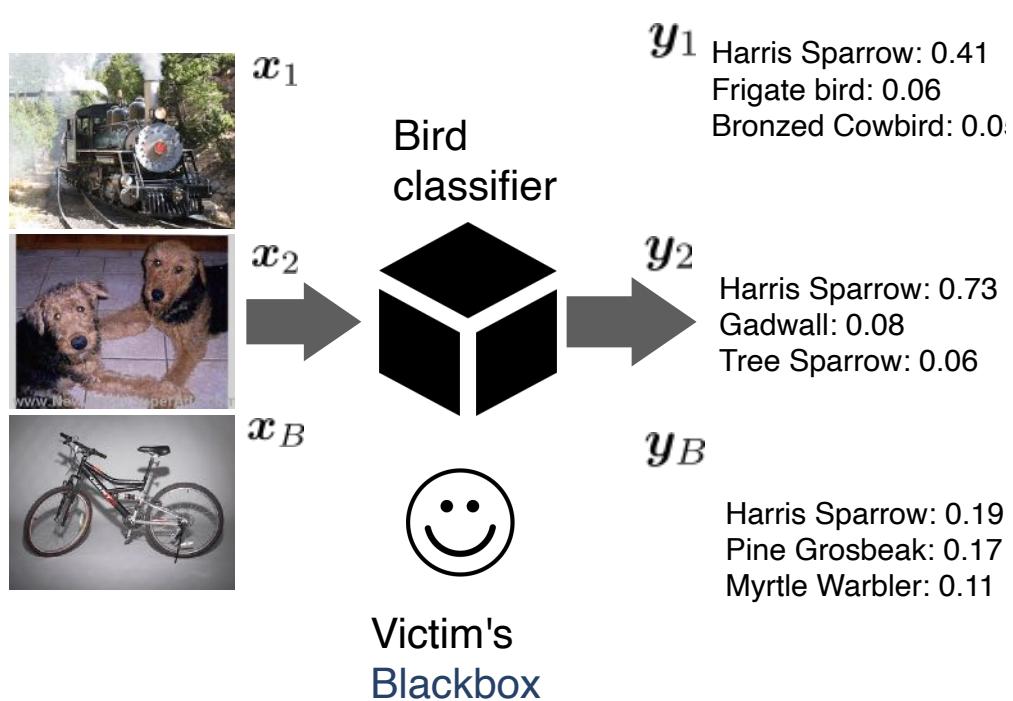
- **Interpretability, Robustness** and **Security** of Deep Learning in Computer Vision
 - ▶ **Inherently Interpretable** Deep Neural networks — CVPR'21, CVPR'22
 - ▶ **Robustness** of Deep Models:
Bright and Dark Side of Scene Context — NeurIPS'18, CVPR'19, ECCV'20
 - ▶ **Security** of Deep Models
Reverse Engineering and **Stealing** of Deep Models — ICLR'18, CVPR'19, ICLR'20

Providing ML Models is a Business Model

- **Input in, prediction out.** Ask \$ per query.
 - ▶ ML models are **black boxes** !
 - ▶ not shared: **architecture, parameters, hyperparameter** details (IPs)
- **Research question:**
 - ▶ can an adversary **steal the functionality of the model** ?

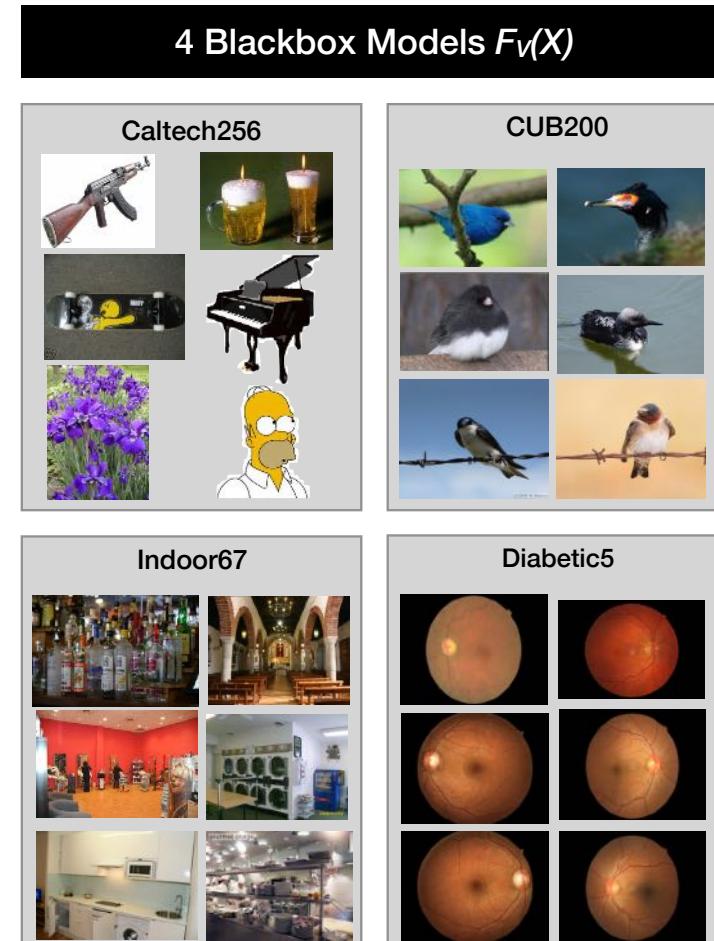


Functionality Stealing: Knock-Off Nets



Transfer Set Construction: $x_i \stackrel{\pi}{\sim} P_A(X)$

- Simple method: $\pi = \text{random}$
 - ▶ sample images randomly (without replacement)
 - ▶ prone to querying irrelevant images



Can we Learn with $\pi = \text{Random?}$ Yes!

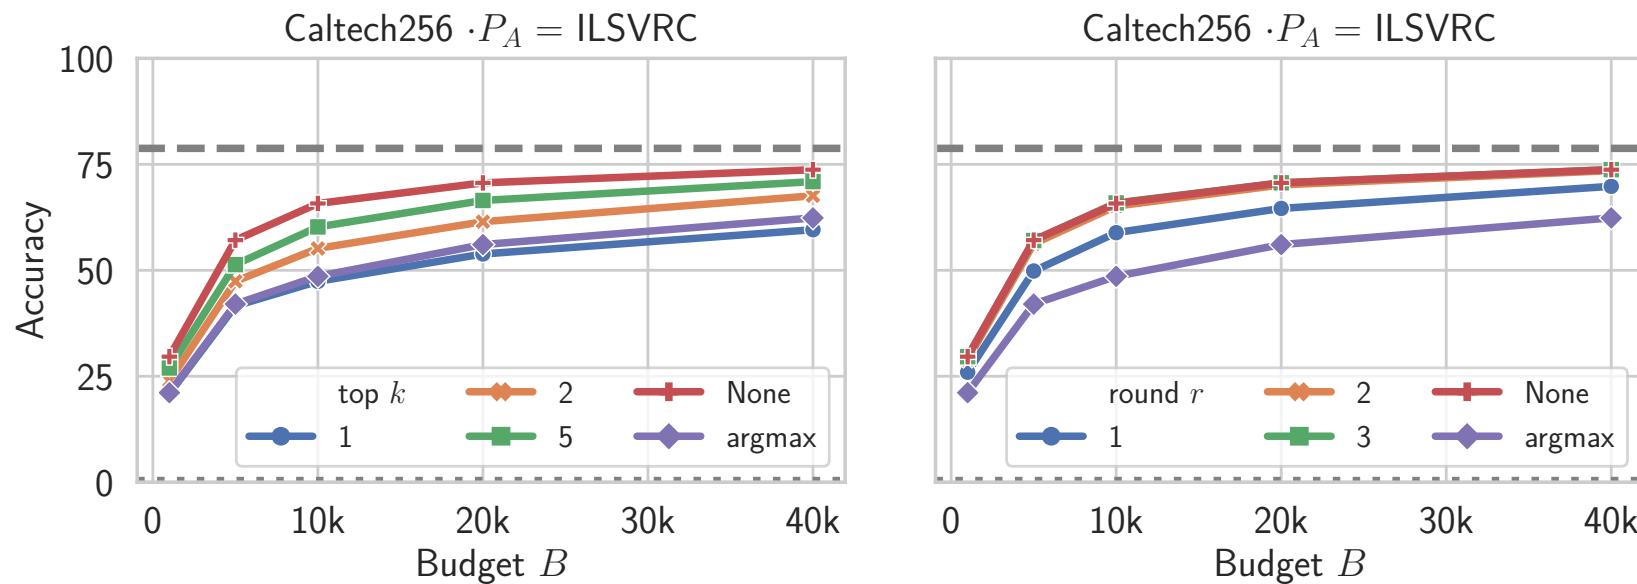
		random			
P_A		Caltech256	CUBS200	Indoor67	Diabetic5
Closed	$P_V(F_V)$	78.8 (1×)	76.5 (1×)	74.9 (1×)	58.1 (1×)
	$P_V(\text{KD})$	82.6 (1.05×)	70.3 (0.92×)	74.4 (0.99×)	54.3 (0.93×)
Closed	D^2	76.6 (0.97×)	68.3 (0.89×)	68.3 (0.91×)	48.9 (0.84×)
Open	ILSVRC	75.4 (0.96×)	68.0 (0.89×)	66.5 (0.89×)	47.7 (0.82×)
	OpenImg	73.6 (0.93×)	65.6 (0.86×)	69.9 (0.93×)	47.0 (0.81×)

accuracy(victim blackbox)

accuracy(knockoff)

$\Rightarrow > 0.81 \times$ accuracy of blackbox recovered

Learning with Less Information? Yes!



⇒ Robust to various **passive defense mechanisms**:
e.g. argmax, top-k, rounding, ...

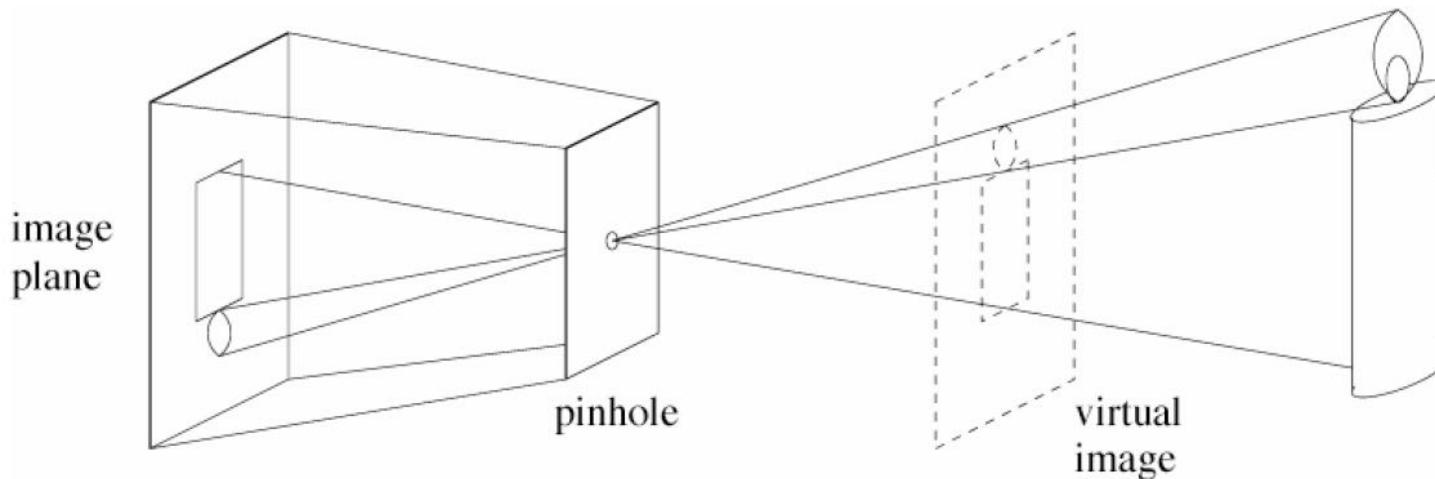


Basic Concepts and Terminology

Computer Vision vs. Computer Graphics

Pinhole Camera (Model)

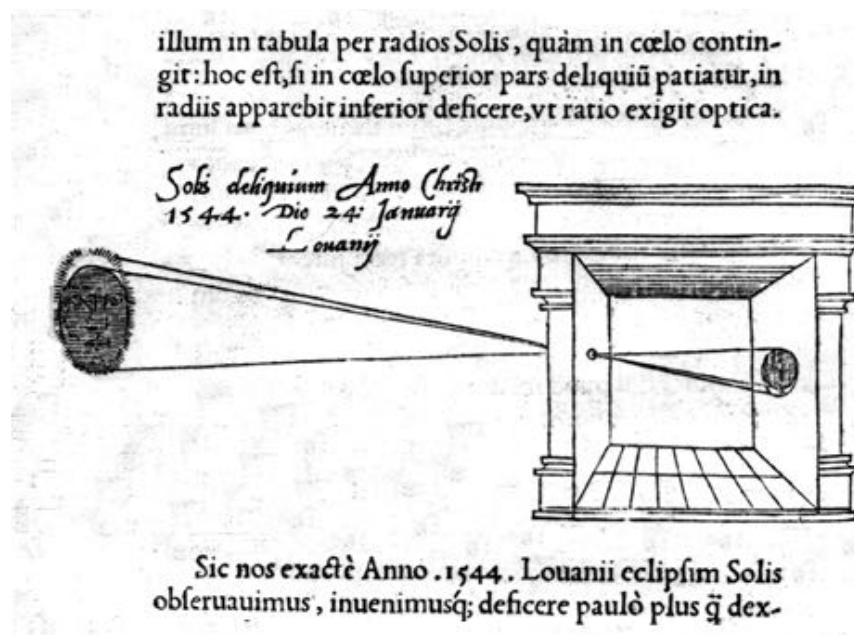
- (simple) standard and abstract model today
 - ▶ box with a small hole in it



Camera Obscura

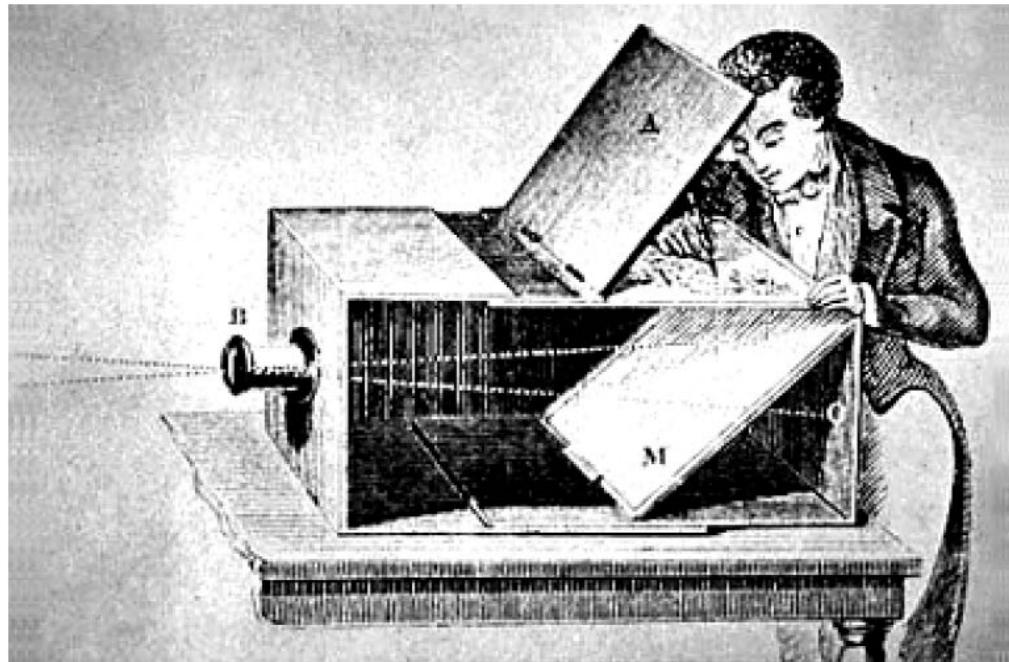
- around 1519, Leonardo da Vinci (1452 - 1519)
 - ▶ http://www.acmi.net.au/AIC/CAMERA_OBSCURA.html

▶ “when images of illuminated objects ... penetrate through a small hole into a very dark room ... you will see [on the opposite wall] these objects in their proper form and color, reduced in size ... in a reversed position owing to the intersection of the rays”



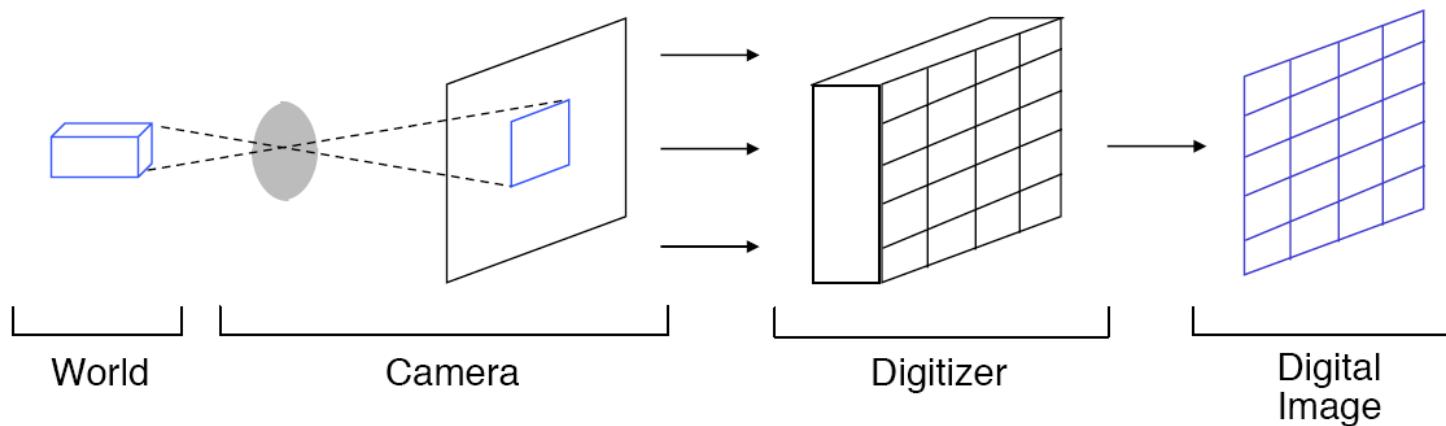
Principle of pinhole....

- ...used by artists
 - ▶ (e.g. Vermeer
17th century,
dutch)
- and scientists



Digital Images

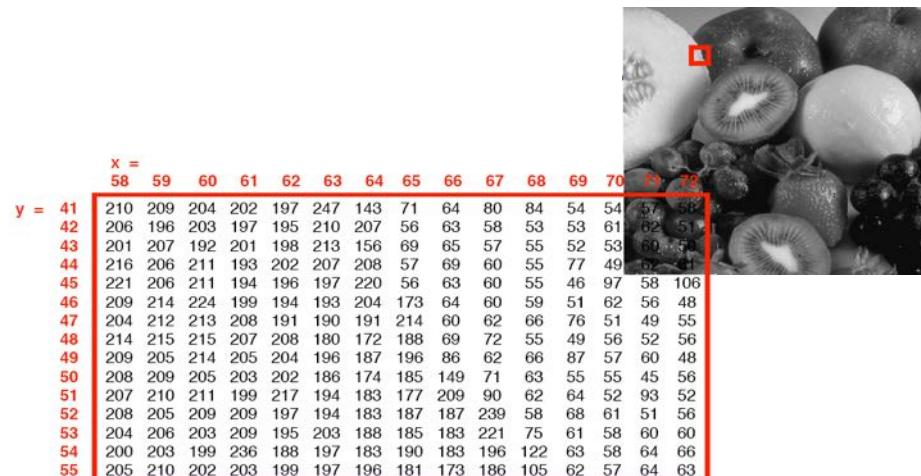
- Imaging Process:
 - ▶ (pinhole) camera model
 - ▶ digitizer to obtain digital image



(Grayscale) Image

- ‘Goals’ of Computer Vision
 - ▶ how can we recognize fruits from an array of (gray-scale) numbers?
 - ▶ how can we perceive depth from an array of (gray-scale) numbers?
 - ▶ ...
- computer vision = the problem of ‘inverse graphics’ ...?

- ‘Goals’ of Graphics
 - ▶ how can we generate an array of (gray-scale) numbers that looks like fruits?
 - ▶ how can we generate an array of (gray-scale) numbers so that the human observer perceives depth?
 - ▶ ...





Case Study

... object recognition

Case Study: Computer Vision & Object Recognition

- how do you recognize
 - ▶ the banana?
 - ▶ the glass?
 - ▶ the towel?
- how can we make computers to do this?
- ill posed problem:
 - ▶ missing data
 - ▶ ambiguities
 - ▶ multiple possible explanations



Image Edges: What are edges? Where do they come from?

- Edges are changes in pixel brightness

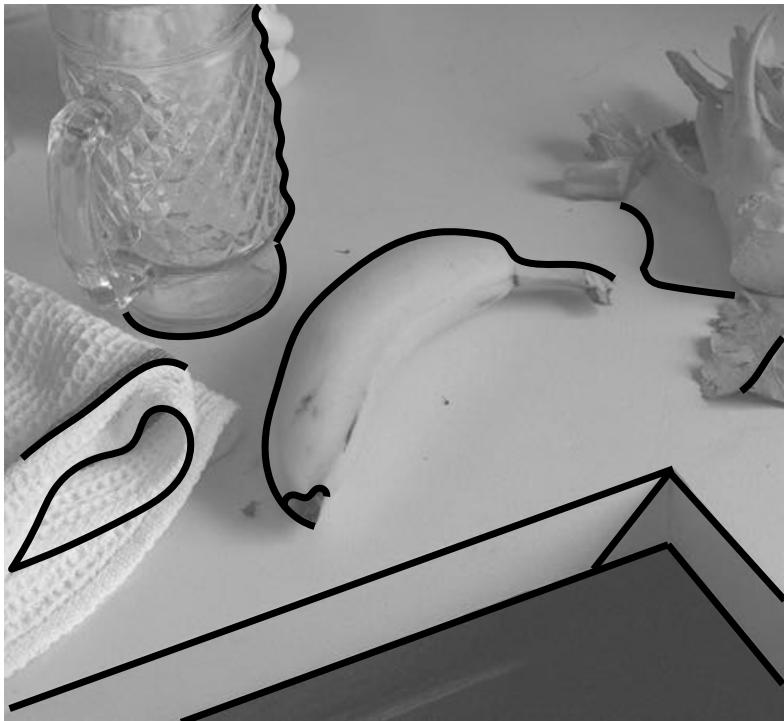
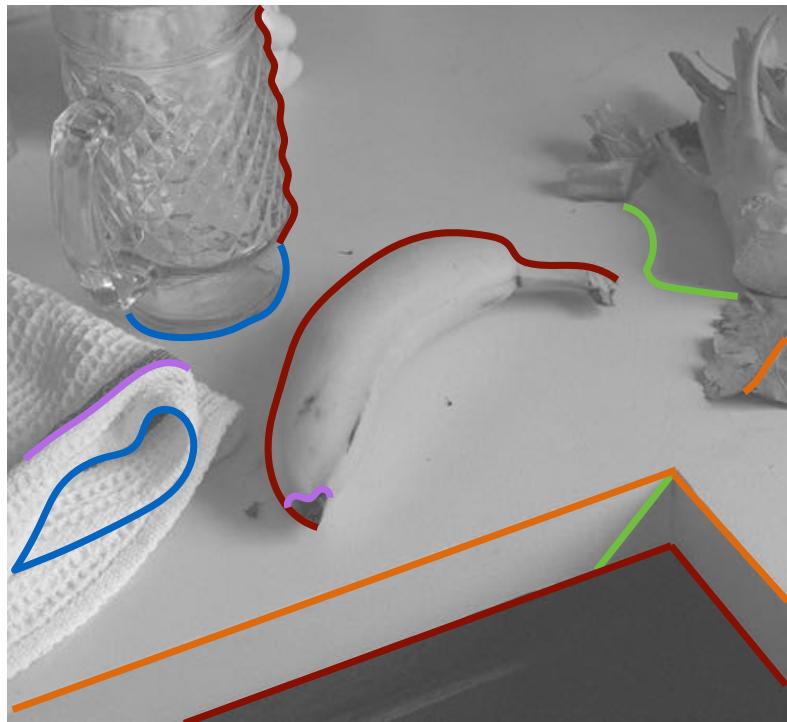


Image Edges: What are edges? Where do they come from?



- Edges are changes in pixel brightness
 - ▶ **Foreground/Background Boundaries**
 - ▶ **Object-Object-Boundaries**
 - ▶ **Shadow Edges**
 - ▶ **Changes in Albedo or Texture**
 - ▶ **Changes in Surface Normals**

Line Drawings: Good Starting Point for Recognition?



MASSACHUSETTS INSTITUTE OF TECHNOLOGY
PROJECT MAC

Artificial Intelligence Group
Vision Memo. No. 100.

July 7, 1966

THE SUMMER VISION PROJECT

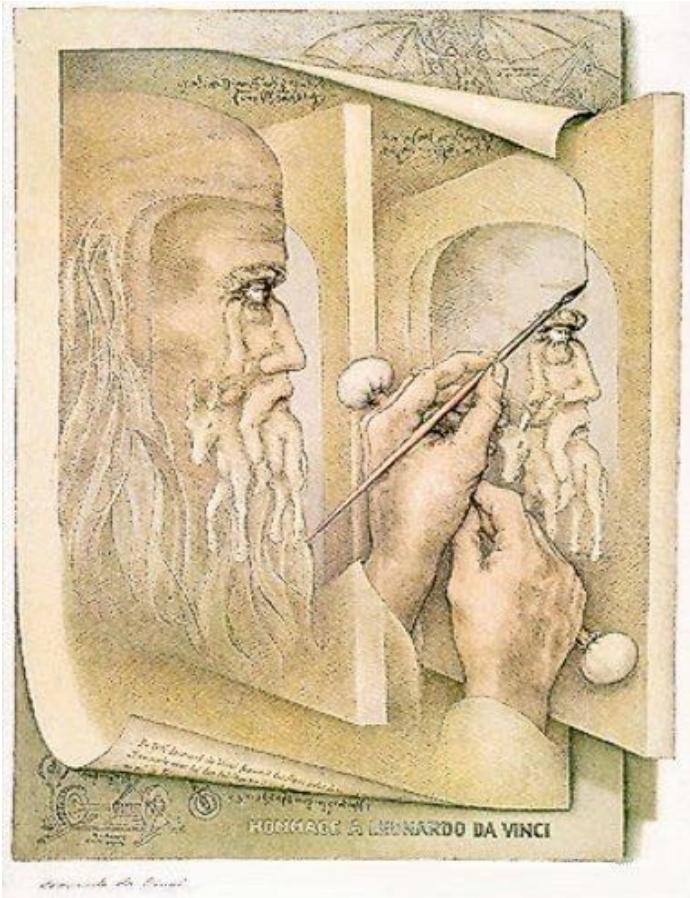
Seymour Papert

The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

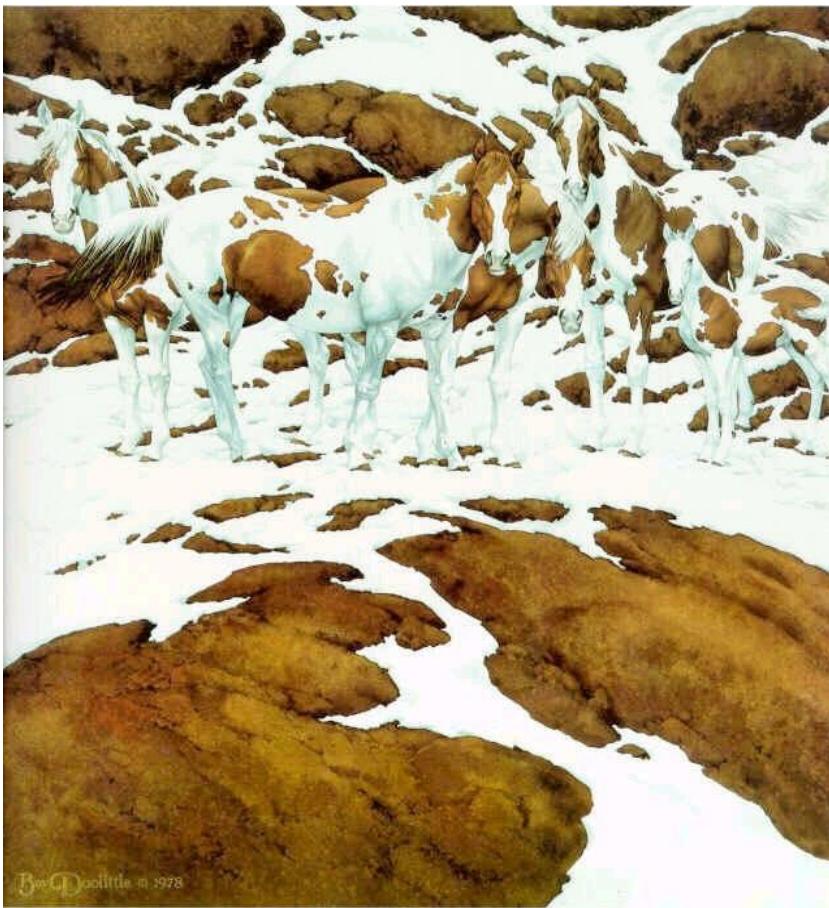
slide credit: Fei-Fei, Justin Johnson, Serena Yeung



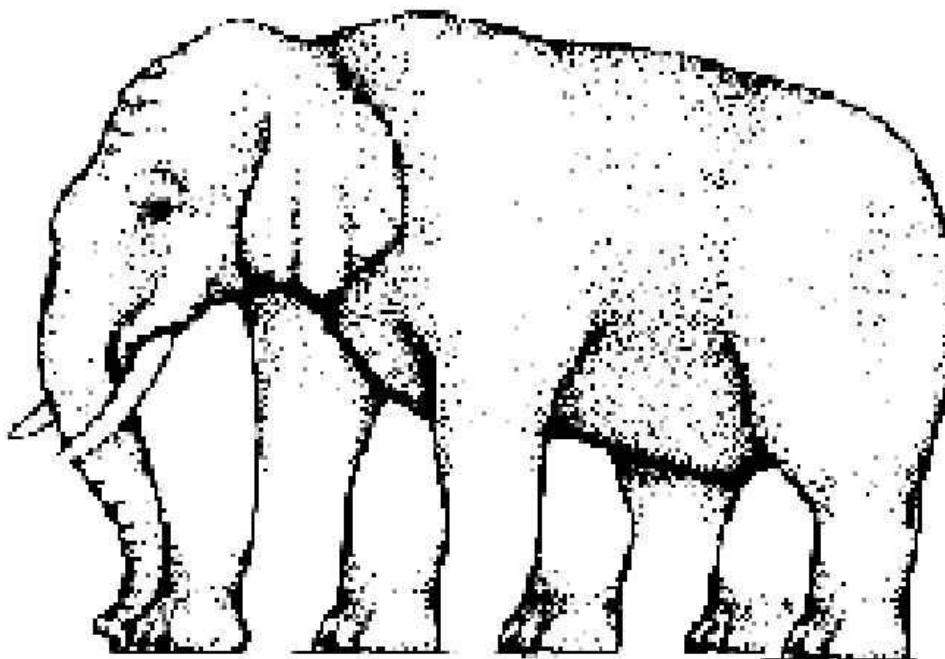
Complexity of Recognition



Complexity of Recognition



Complexity of Recognition

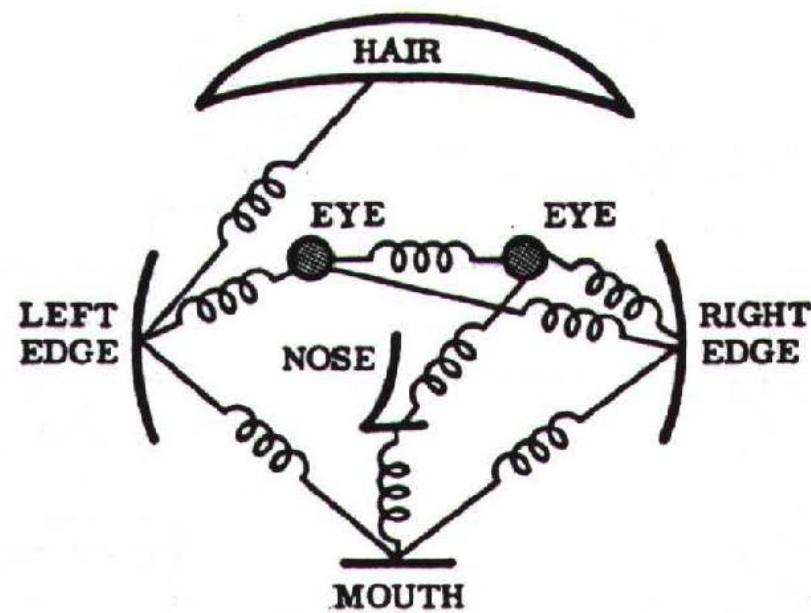


Complexity of Recognition



Class of Models: Pictorial Structure

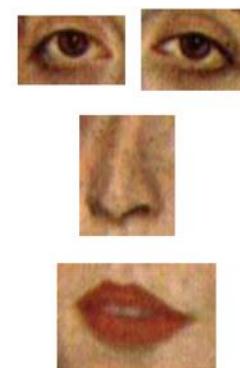
- Fischler & Elschlager 1973
- Model has two components
 - ▶ parts
(2D image fragments)
 - ▶ structure
(configuration of parts)



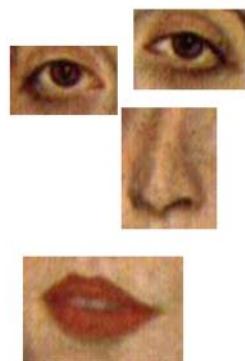
Deformations



A



B



C

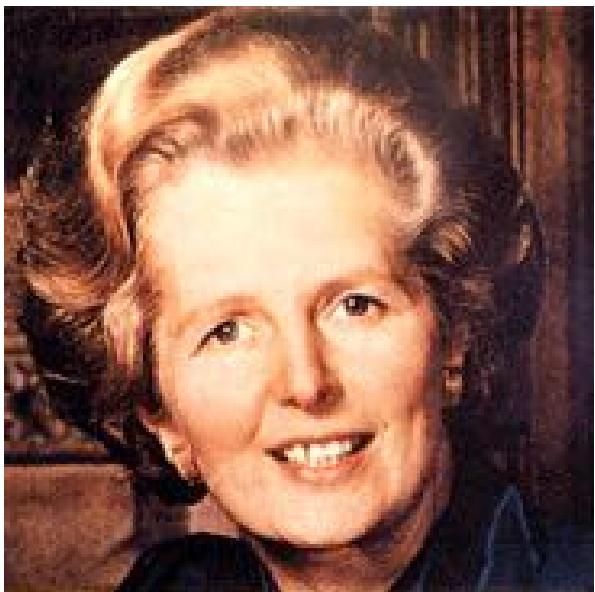


D

Clutter



Example





Recognition, Localization, and Segmentation

a few terms

... let's briefly define what we mean by that

Object Recognition: First part of this Computer Vision class

- Different Types of Recognition Problems:
 - ▶ Object **Identification**
 - recognize your pencil, your dog, your car
 - ▶ Object **Classification**
 - recognize any pencil, any dog, any car
 - also called: generic object recognition, object categorization, ...
- Recognition and
 - ▶ **Segmentation**: separate pixels belonging to the foreground (object) and the background
 - ▶ **Localization/Detection**: position of the object in the scene, pose estimate (orientation, size/scale, 3D position)

Object Recognition: First part of this Computer Vision class

- Different Types of Recognition Problems:

- ▶ Object **Identification**

- recognize your apple,
your cup, your dog

- ▶ Object **Classification**

- recognize any apple,
any cup, any dog
 - also called:
generic object recognition,
object categorization, ...
 - typical definition:
'basic level category'



Which Level is right for Object Classes?

- Basic-Level Categories
 - ▶ the highest level at which category members have **similar perceived shape**
 - ▶ the highest level at which a **single mental image** can reflect the entire category
 - ▶ the highest level at which a person uses similar **motor actions** to interact with category members
 - ▶ the level at which human subjects are usually **fastest** at identifying category members
 - ▶ the first level named and understood by **children**

 - ▶ (while the definition of basic-level categories depends on culture there exist a remarkable consistency across cultures...)
- Most recent work in object recognition has focused on this problem
 - ▶ we will discuss several of the most successful methods in the lecture :-)

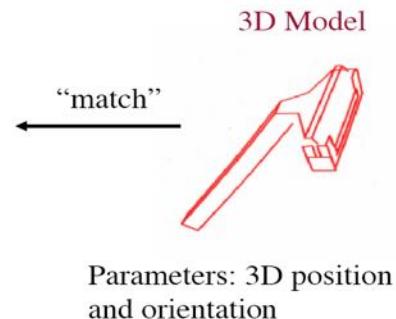
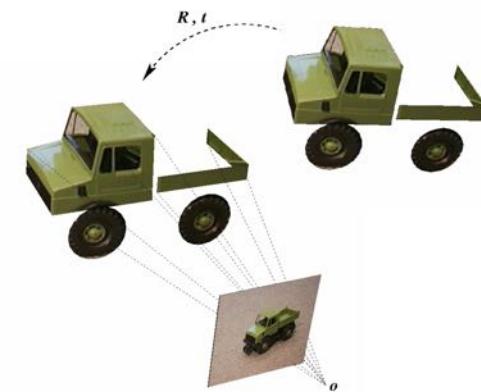
Object Recognition & Segmentation

- Recognition and
 - ▶ **Segmentation**: separate pixels belonging to the foreground (object) and the background



Object Recognition & Localization

- Recognition and
 - ▶ **Localization in 3D**: to position the object in the scene, estimate the object's pose (orientation, size/scale, 3D position)
 - ▶ Example from David Lowe:



Object Recognition

- Different Types of Recognition Problems:
 - ▶ Object **Identification**
 - recognize your pencil, your dog, your car
 - ▶ Object **Classification**
 - recognize any pencil, any dog, any car
 - also called: generic object recognition, object categorization, ...
- Recognition and
 - ▶ **Segmentation**: separate pixels belonging to the foreground (object) and the background
 - ▶ **Localization**: position the object in the scene, estimate pose of the object (orientation, size/scale, 3D position)

Goals of today's lecture

- First intuitions about
 - ▶ What is computer vision?
 - ▶ What does it mean to see and how do we (as humans) do it?
 - ▶ How can we make this computational?
- Applications & Appetizers
- Case Study
 - ▶ Object Recognition — intuition from human vision...
- Image Classification
 - ▶ using data-driven approaches (machine learning)
 - ▶ K-nearest neighbor classifier

Image Classification: a core task in computer vision



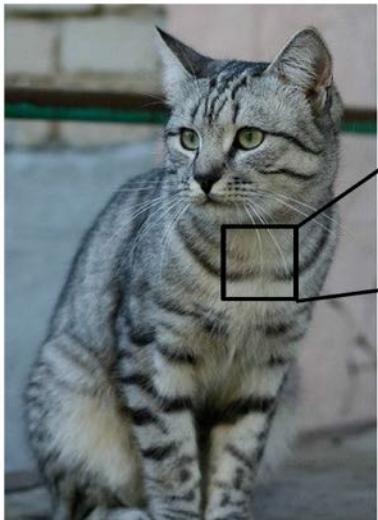
This image by Nikita is
licensed under CC-BY 2.0

(assume given set of discrete labels)
{dog, cat, truck, plane, ...}



Image Classification

The Problem: Semantic Gap



```
[[105 112 108 111 104 99 106 99 96 103 112 119 104 87 93 87]
 [ 91 98 102 106 104 79 98 103 99 105 123 136 110 105 94 85]
 [ 76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]
 [ 99 81 81 93 128 131 127 108 95 98 102 99 96 93 101 94]
 [106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]
 [114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]
 [133 137 147 103 65 81 88 65 52 54 74 84 102 93 85 62]
 [128 137 144 148 109 95 86 70 62 65 63 63 60 73 86 101]
 [125 129 140 137 119 121 117 94 65 79 80 65 54 64 72 98]
 [127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
 [115 114 109 123 158 148 131 118 113 109 100 92 74 65 72 78]
 [ 89 93 90 97 108 147 131 118 113 114 113 109 106 65 77 80]
 [ 63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]
 [ 62 65 82 89 78 71 88 101 124 126 119 101 107 114 131 119]
 [ 63 65 75 88 89 71 62 81 120 138 135 105 81 68 110 118]
 [ 87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]
 [118 97 82 86 117 123 116 66 41 51 95 53 69 65 102 107]
 [164 146 112 88 82 128 124 104 76 48 45 66 68 101 102 109]
 [157 170 157 128 93 86 114 132 112 57 69 55 70 82 69 94]
 [130 128 134 161 139 108 109 118 121 134 114 87 65 53 69 86]
 [128 112 96 117 158 144 128 115 104 107 102 93 87 81 72 79]
 [123 107 96 86 83 112 153 149 122 109 104 75 80 107 112 99]
 [122 121 102 88 82 86 94 117 145 148 153 102 58 78 92 107]
 [122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]]
```

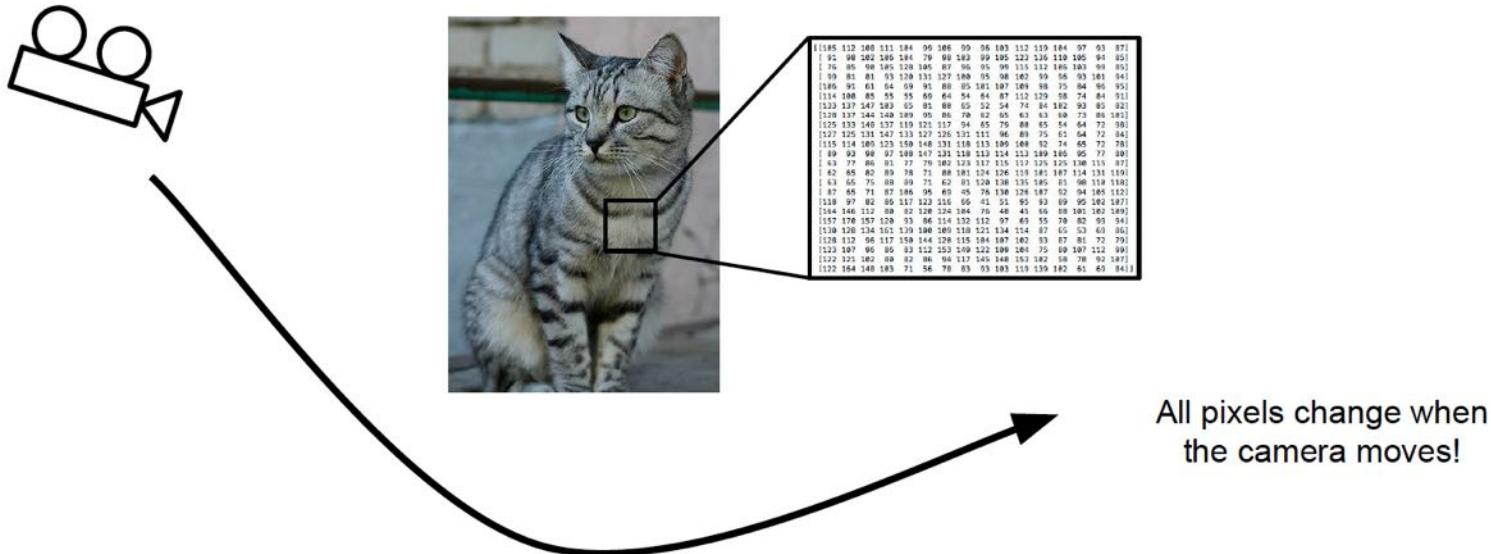
What the computer sees

An image is just a big grid of numbers between [0, 255]:

e.g. 800 x 600 x 3
(3 channels RGB)

Image Classification

Challenges: Viewpoint variation



This image by Nikita is
licensed under CC-BY 2.0



Image Classification

Challenges: Illumination



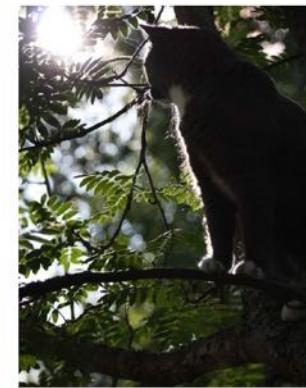
[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

Image Classification

Challenges: Deformation



This image by [Umberto Salvagnin](#)
is licensed under [CC-BY 2.0](#)



This image by [Umberto Salvagnin](#)
is licensed under [CC-BY 2.0](#)



This image by [sare bear](#) is
licensed under [CC-BY 2.0](#)



This image by [Tom ThaIs](#)
is licensed under [CC-BY 2.0](#)

Image Classification

Challenges: Occlusion



[This image](#) is CC0 1.0 public domain



[This image](#) is CC0 1.0 public domain



[This image](#) by [jonsson](#) is licensed under CC-BY 2.0

Image Classification

Challenges: Background Clutter



[This image is CC0 1.0 public domain](#)



[This image is CC0 1.0 public domain](#)

Image Classification

Challenges: Intraclass variation



This image is CC0 1.0 public domain

Image Classification

An image classifier

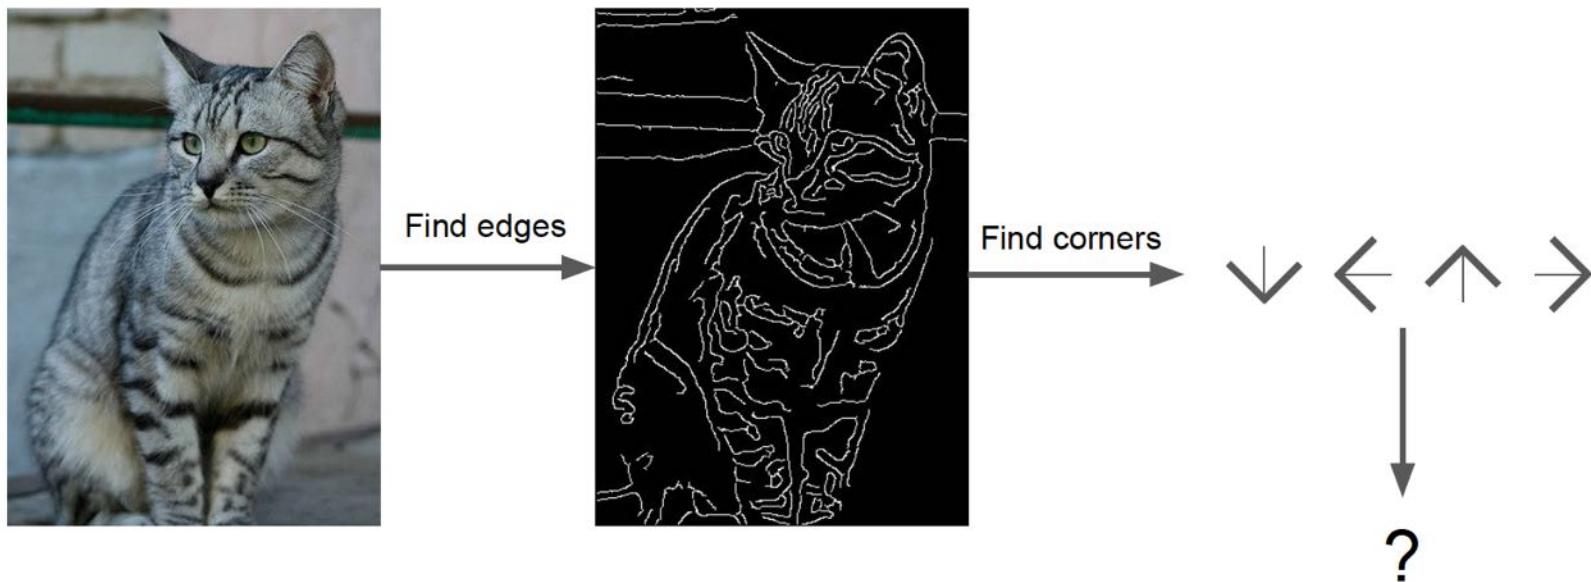
```
def classify_image(image):  
    # Some magic here?  
    return class_label
```

Unlike e.g. sorting a list of numbers,

no obvious way to hard-code the algorithm for
recognizing a cat, or other classes.

Image Classification

Attempts have been made



John Canny, "A Computational Approach to Edge Detection", IEEE TPAMI 1986

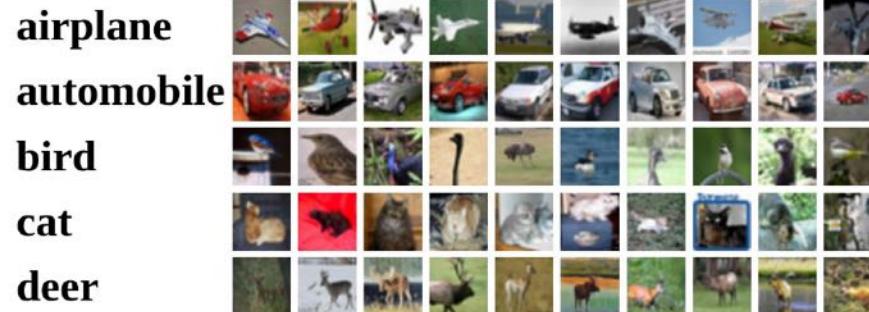
Machine Learning: Data-Driven Approach

1. Collect a dataset of images and labels
2. Use Machine Learning to train a classifier
3. Evaluate the classifier on new images

Example training set

```
def train(images, labels):
    # Machine learning!
    return model

def predict(model, test_images):
    # Use model to predict labels
    return test_labels
```



First Classifier: Nearest Neighbor

```
def train(images, labels):  
    # Machine learning!  
    return model
```

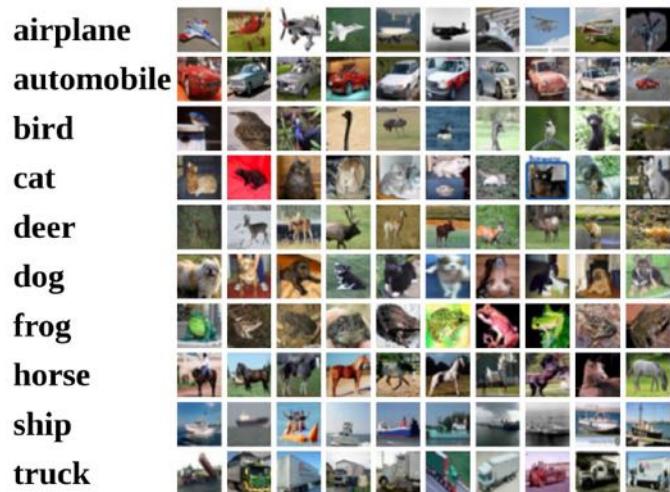
→ Memorize all
data and labels

```
def predict(model, test_images):  
    # Use model to predict labels  
    return test_labels
```

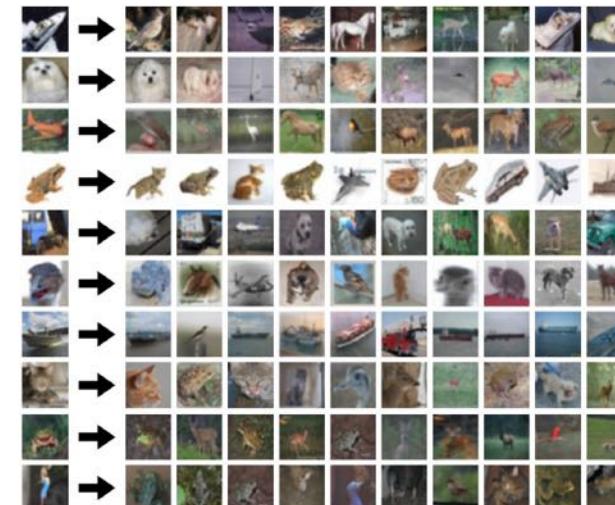
→ Predict the label
of the most similar
training image

Example Dataset: CIFAR10

10 classes
50,000 training images
10,000 testing images



Test images and nearest neighbors



Alex Krizhevsky, "Learning Multiple Layers of Features from Tiny Images", Technical Report, 2009.

First Classifier: Nearest Neighbor Classifier

Distance Metric to compare images

L1 distance:

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$

test image				training image				pixel-wise absolute value differences			
56	32	10	18	10	20	24	17	46	12	14	1
90	23	128	133	8	10	89	100	82	13	39	33
24	26	178	200	12	16	178	170	12	10	0	30
2	0	255	220	4	32	233	112	2	32	22	108

-

=

add → 456

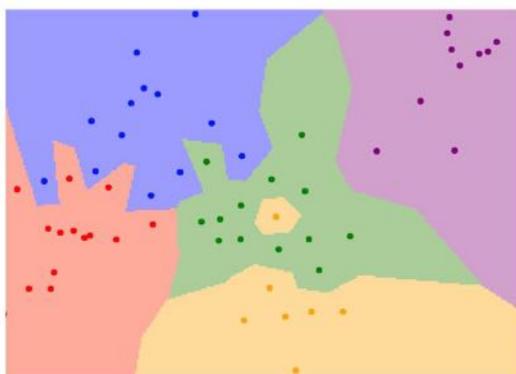
First Classifier: Nearest Neighbor Classifier

What does this look like?

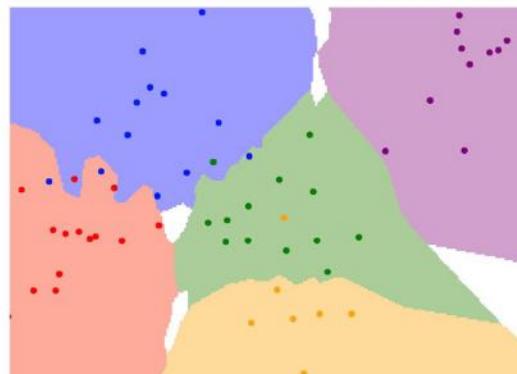


K-Nearest Neighbor Classifier

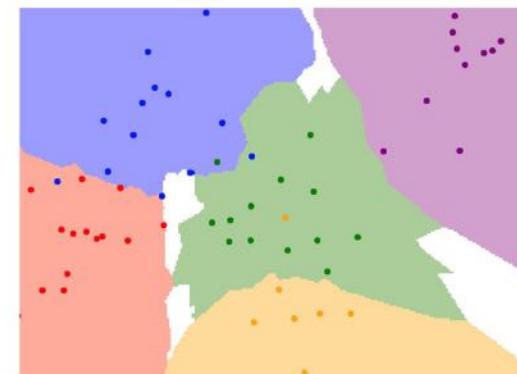
Instead of copying label from nearest neighbor,
take **majority vote** from K closest points



$K = 1$



$K = 3$



$K = 5$

K-Nearest Neighbor Classifier

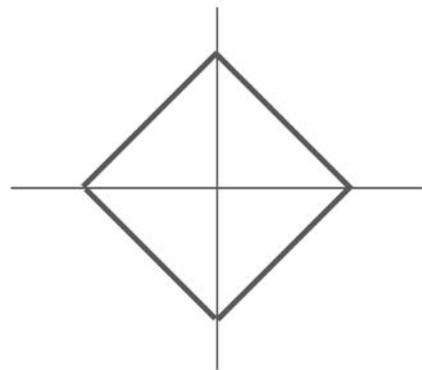
What does this look like?



K-Nearest Neighbor Classifier: Distance Metric

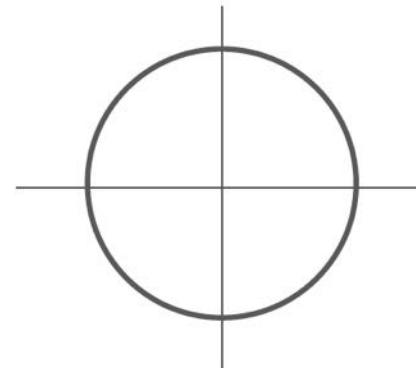
L1 (Manhattan) distance

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



L2 (Euclidean) distance

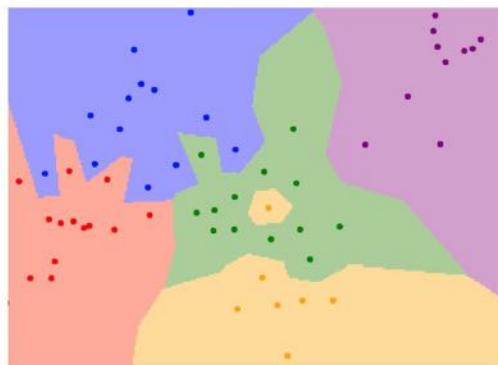
$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



K-Nearest Neighbor Classifier: Distance Metric

L1 (Manhattan) distance

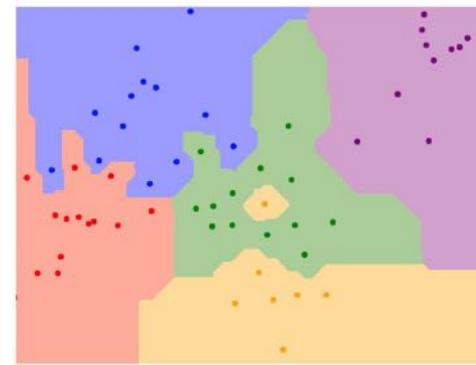
$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



$K = 1$

L2 (Euclidean) distance

$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



$K = 1$

Hyperparameters

What is the best value of **k** to use?

What is the best **distance** to use?

These are **hyperparameters**: choices about the algorithm that we set rather than learn

Very problem-dependent.

Must try them all out and see what works best.

Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: K = 1 always works perfectly on training data

Your Dataset

Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data

train

test

Idea #3: Split data into **train**, **val**, and **test**; choose hyperparameters on val and evaluate on test

Better!

train

validation

test

Setting Hyperparameters

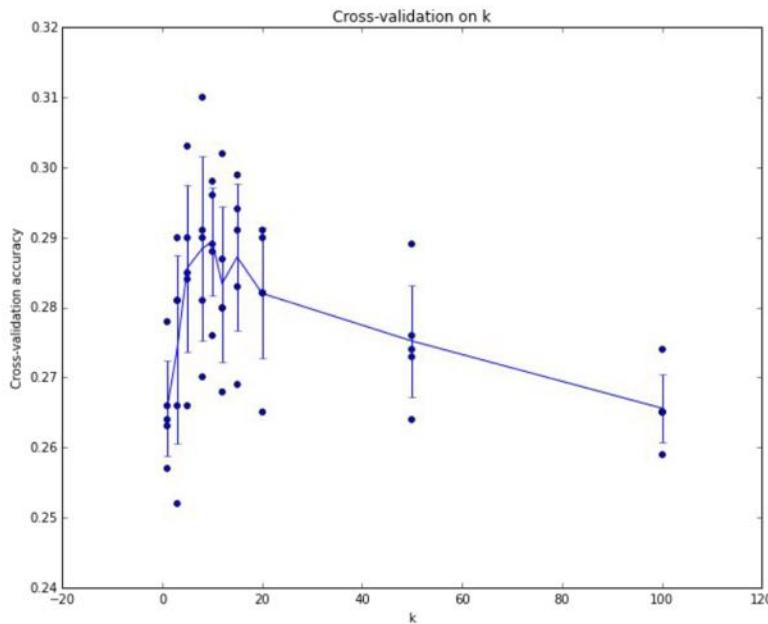
Your Dataset

Idea #4: Cross-Validation: Split data into **folds**,
try each fold as validation and average the results

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

Useful for small datasets, but not used too frequently in deep learning

Setting Hyperparameters



Example of
5-fold cross-validation
for the value of **k**.

Each point: single
outcome.

The line goes
through the mean, bars
indicated standard
deviation

(Seems that $k \approx 7$ works best
for this data)

Nearest Neighbor - not used for images :-)

- Very slow at test time
- Distance metrics on pixels are not informative



Original image is
CC0 public domain

(all 3 images have same L2 distance to the one on the left)

K-Nearest Neighbors: Summary

In **Image classification** we start with a **training set** of images and labels, and must predict labels on the **test set**

The **K-Nearest Neighbors** classifier predicts labels based on nearest training examples

Distance metric and K are **hyperparameters**

Choose hyperparameters using the **validation set**; only run on the test set once at the very end!

Linear Classification

